Stress-Testing AI:
How to Evaluate Whether Your AI Tools Are Working

Jonathan Soma
js4571@columbia.edu @dangerscarf
Columbia University

Two definitions of “working”

- **Serving a true user need** vs **the tech is working**
- Example:
  - “We made an AI chatbot to answer questions about elections”
  - **Does the reader want to engage with it?** Not for today!
  - **Does it give accurate responses?** This is our session!

Fear not, failure is normal

- Every failure just becomes one more thing to test for
- Don’t cry because it failed, smile because it happened :/
- The stakes are lower than you think, I swear, every newsroom has screwed up
- But you do have an obligation to not screw up!

AI translation now 'good enou

pressgazette.co.uk/publishers/h...

In this section

The AI-translated text is not edited by humans because
said, the "workflow is so tight" on Espresso which upda
around 20 times a day.

"There is no natural thing where we can say 'okay, now
everything is done. Let's translate, and let's look at the
translations and make sure they're perfect'. That doesn
The only thing we can do is, if it's really embarrassing,
take it down and the next version in 20 minutes will be

One embarrassing exa
turned German Chanc

But Siegele said a Fren
"I don't read English. Th
Economist without having to put it into Google Transla
get bad translations."

The Economist's AI-translated soci
videos

The Economist simultaneously launched AI-translated
on its social platforms in the same four languages.

The videos are all a maximum of 90 seconds meaning i

Washington Post's AI-genera

semafor.com/article/12/1...

Events | Email Briefings

8:54 AM FRIDAY JANUARY 9, 2026

D.C. BXL LAGOS RIYADH

SEMAFO

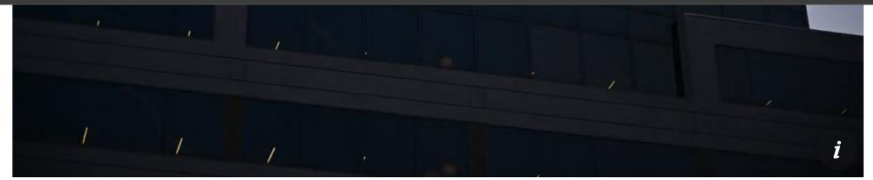
Intelligence for the New World Economy

Exclusive / Washington Post's AI-ge
podcasts rife with



LA Times reportedly removes

theguardian.com/us-news/2...



Los Angeles

This article is more than 1 year old

LA Times reportedly removes new AI tool from story after it downplayed KKK

A day after launch of contentious feature, its bias meter said some accounts see Klan as 'responding to societal changes'

The Los Angeles Times has reportedly removed its new artificial intelligence (AI) tool from one of its articles, just a day after launching the feature.

The publication unveiled a new AI-driven feature called "Insights" this week, designed to accompany articles that express a particular viewpoint or are "written from a personal perspective", such as opinion pieces, commentary and reviews.

The purpose of the tool is to help readers understand where the "views expressed may fall on the political spectrum" and to provide annotated summaries "of the ideas expressed in the piece along with different views on the topic from a variety of sources", according to the LA Times website.

This is the catch

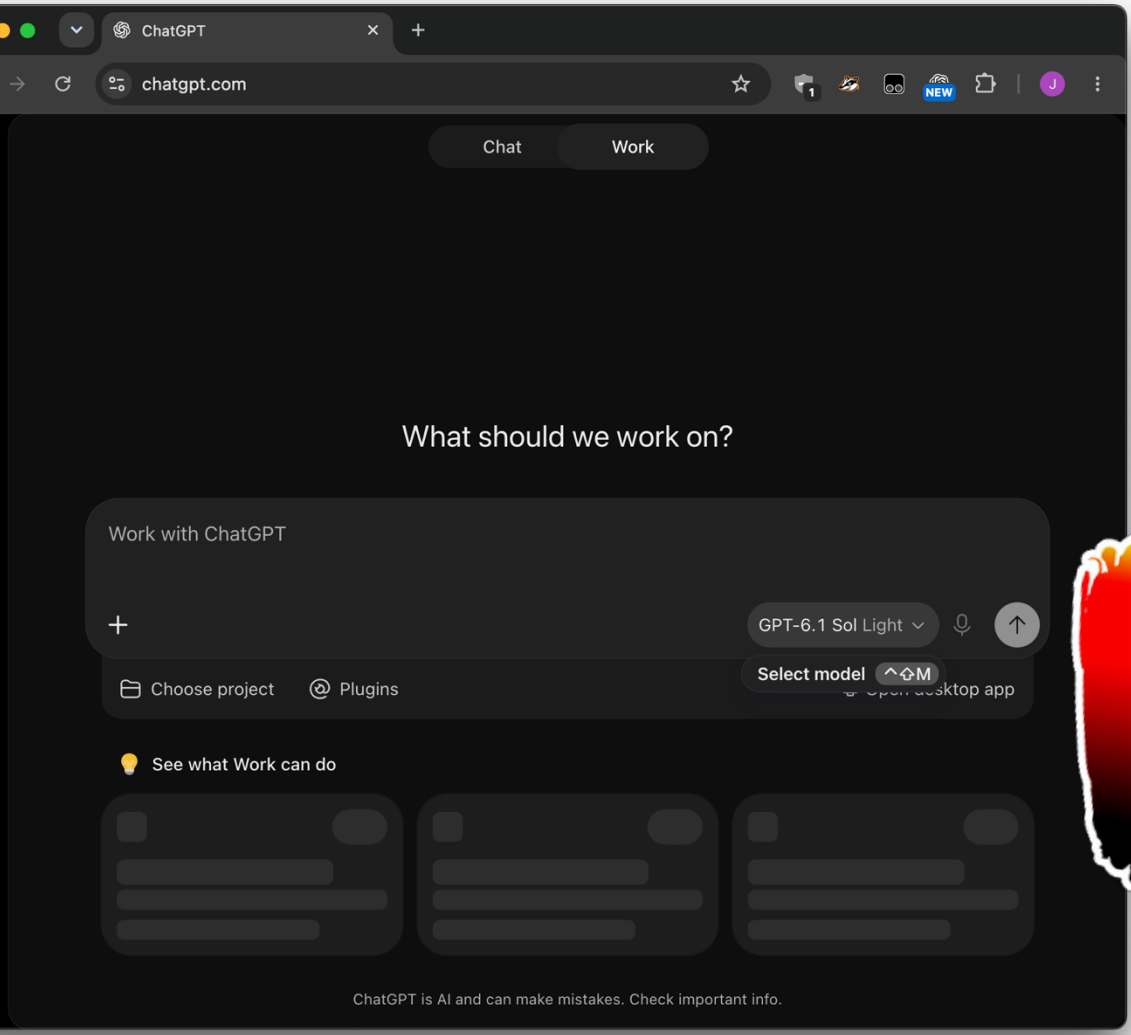
Support us

Maybe later

they all still exist!!!!!

What we are doing today

- Checking the outputs of your AI
- Checking the process inside of your AI
- What the tools look like, existing and bespoke
- Minimize cost while maximizing performance
- What to do with failures?
- Theory, then tools
- And finally: a couple case studies



VS

```python
from openai import OpenAI
from somewhere import openfec_data, exa_search

client = OpenAI()

# Imaginary helper converts Python functions
into API tool definitions.
response = client.responses.create(
    model="gpt-5",
    input="Research {name}'s businesses and
political donations.",
    tools=[
        tool(exa_search),
        tool(openfec_data),
    ],
)
```

Our example use case

- We'll look at each step through the lens of a **simple research agent** that delivers a report.
  - Web research
  - Tool use
  - Presentation
- At the end we'll look at a simpler newsletter summarizer and an "ask our archives" chatbot

THE FOLLOWING NAMED PERSON HAS REQUESTED PERMISSION TO ADDRESS THE CITY COUNCIL ON MONDAY, OCTOBER 5, 2026 AT 6:30 P.M. IN THE CITY COUNCIL CHAMBERS, CITY HALL, 45 LYON TERRACE, BRIDGEPORT, CT 06604.

NAME	SUBJECT
John Marshall Lee 30 Beacon Street Bridgeport, CT 06605	Civics – Municipal Governance.

Mr. Lee came forward and made the following statement:

My name is John Marshall Lee. I live at 30 Beacon Street. And for those of you who are on the Council, I've distributed a copy of an article from this week's October 3rd, The Economist, and it's about it's a writer for The Economist magazine wrote, "But there's one

The question became more complicated for voters to become informed. Where did the separate request come from originate? Only the City Clerk appeared at one hearing and one sheriff also appeared to share their thoughts. Who lobbied for the question to become divided. Perhaps you remember the result.

Question 1 on the charter change received yes approval from 75% of those voting. You're going to get the complete results because I just got them from the charter commission after 11 months from the Registrar of Voters office. So, you will hear those next time and we'll pick out the people and where they came from.

John Farrell 751 Norman Street Bridgeport, CT 06605	Handicap signs.

Council President Pro Tem Newton called Mr. Farrell. There was no response.

Alicia A. Carpenter-Cummings Amsterdam Avenue Bridgeport, CT 06606	Ox Brook Flooding and Water 534 drainage on Amsterdam Avenue.

Ms. Carpenter-Cummings came forward and made the following statement:

Good evening, members of the City Council. My name is Alicia Carpenter, the homeowner of 534 Amsterdam Avenue in the 134th district.

§ WHO HELD THE KEYS

What memes failed to disclose, instead, was who held the keys.

The Ledger's analysis identified 4,146 confirmed rug pulls — schemes in which a token's developers drain its liquidity pool and disappear — across the 24-month window. Reported retail losses peaked at $217 million in April 2024.

The pattern echoes the 2022 collapse of the Squid Game token, which cratered to zero overnight and whose developers were never identified.

"I sent $14,000 in USDC to a wallet I found on Telegram," said Daniel Reyes, a 31-year-old delivery driver in Phoenix who lost his savings to a Solana-based meme coin called "ElonRocket" in March 2024. "Twelve hours later the chart was a single red line."

A regression of monthly rug-pull losses against meme-coin market capitalization in the Ledger's dataset showed a correlation coefficient of 0.94, suggesting that fraud scales almost mechanically with hype rather than with prices.

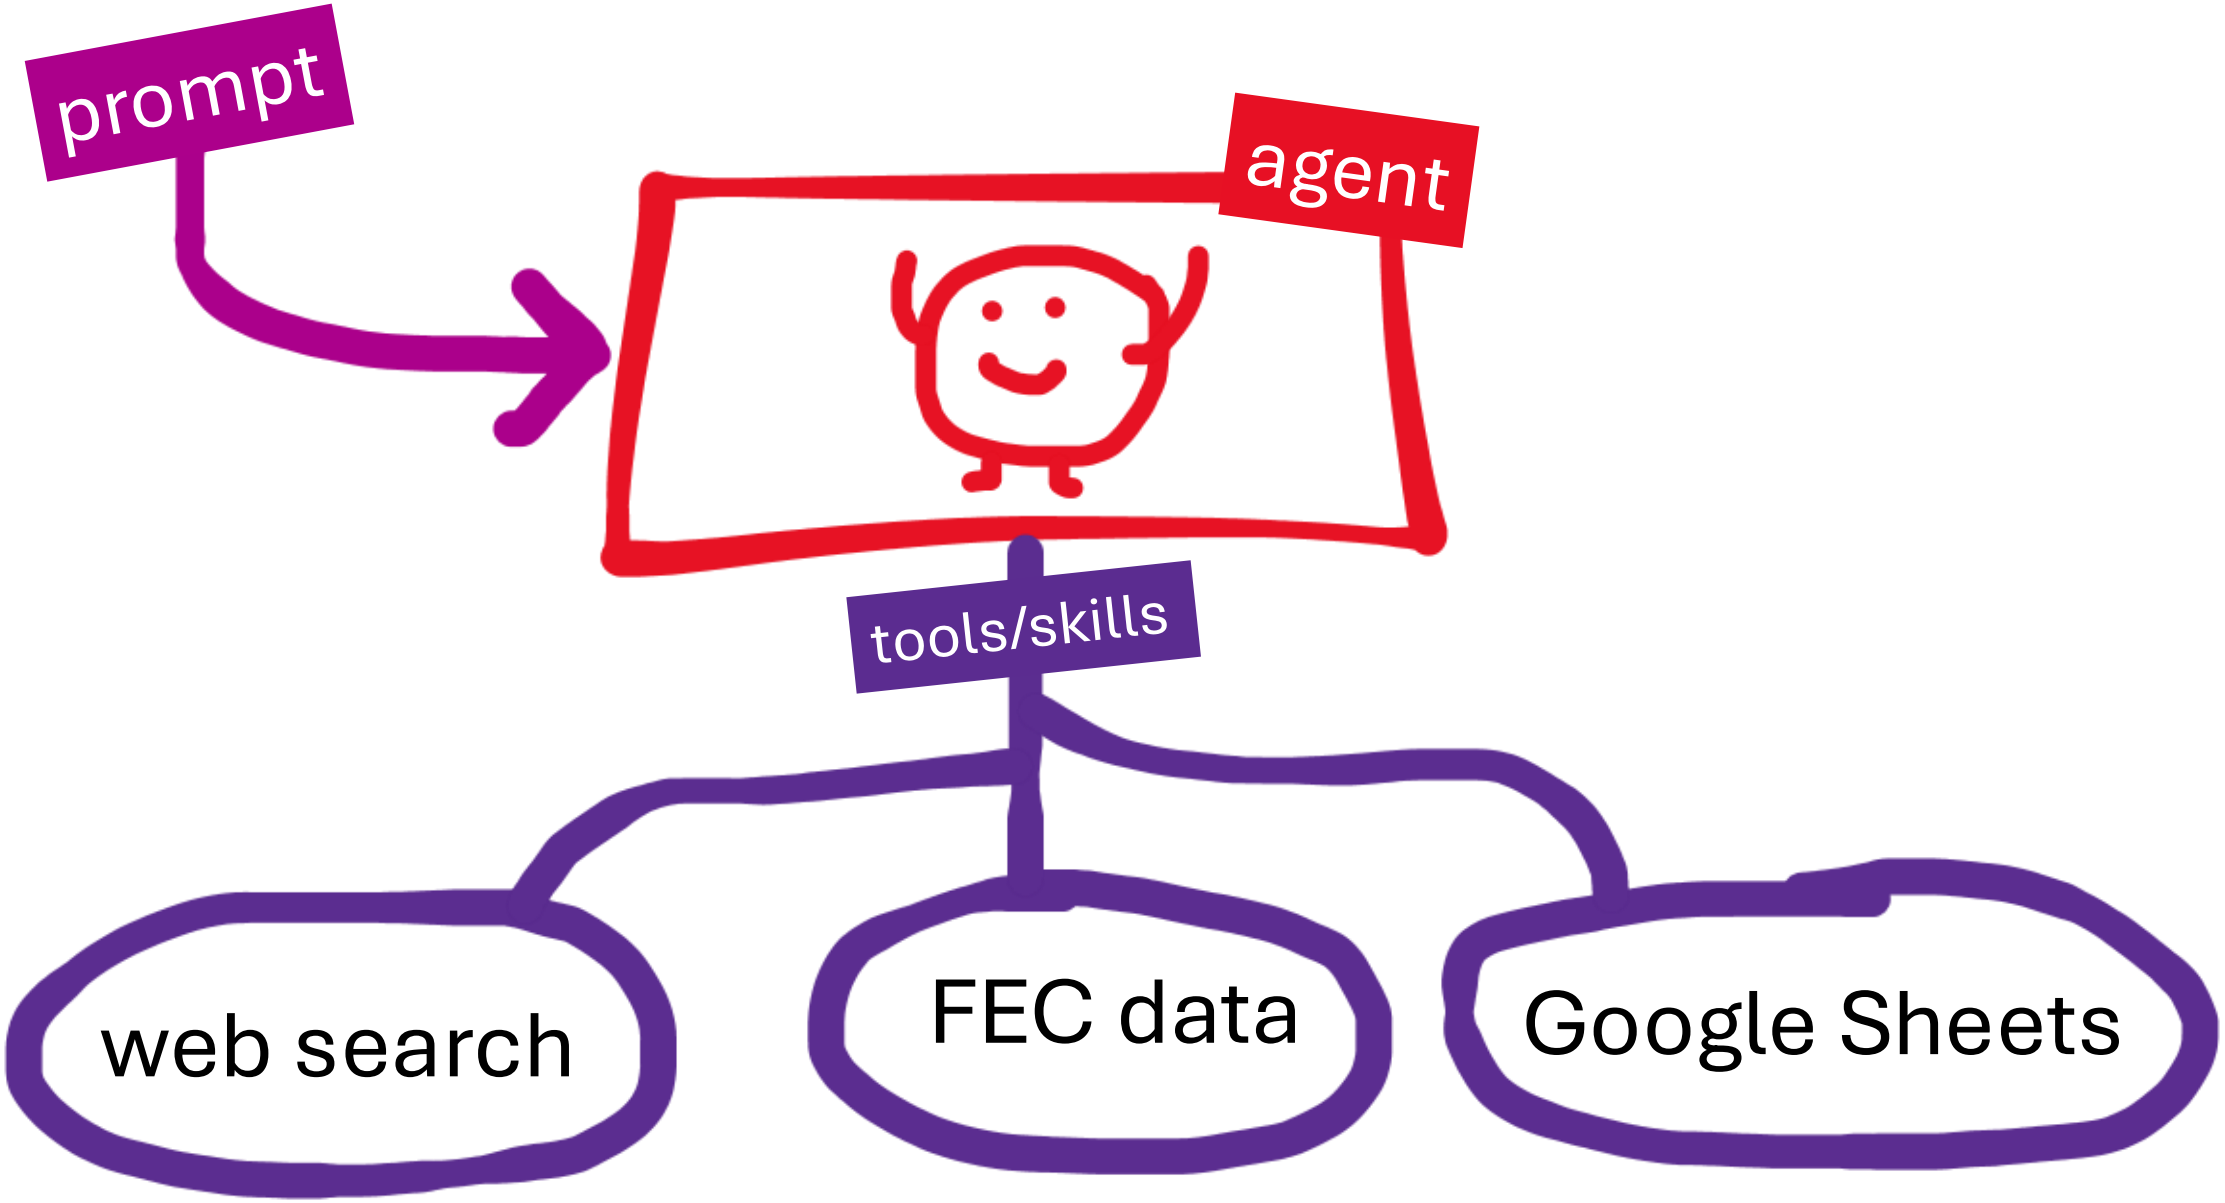
§ A CYCLE, OR A CREST?

Whether the cycle has crested is unclear.

Bitcoin's rally past $93,000 has restored institutional interest in the broader sector. Cathie Wood, the chief executive of ARK Invest, has reiterated her forecast that bitcoin will trade above $1 million within a decade.

But for the retail investors who chased frogs and dogs across blockchains, the lesson may be a familiar one.

STEP ONE:
Get to a good place


prompt
agent
tools/skills
web search
FEC data
Google Sheets

“Look this person up and find a little summary about them and grab their LinkedIn and details on any campaign contributions: {name}”



Elizabeth Uihlein

Co-founder and president of Uline, the Wisconsin-based packaging and shipping supplies company. She and her husband, Richard Uihlein, are major contributors to Republican candidates and political committees. DeSmog +1

LinkedIn

I couldn't verify a personal LinkedIn profile for Elizabeth Uihlein. Searches for both “Elizabeth Uihlein” and “Liz Uihlein,” including searches restricted to LinkedIn profiles, did not surface a credible match.

Search Elizabeth Uihlein on LinkedIn ↗.

Federal campaign contributions

I found federal FEC filings identifying Elizabeth A. Uihlein as a contributor, often listing her employer as Uline and her occupation as CEO, president, or co-founder.

Here are some documented examples:

Contribution date	Recipient	Amount
May 21, 2026	Julia Letlow for Louisiana	$3,500 × 2
October 14, 2025	Federal candidate committee (filing excerpt)	$3,500
June 10, 2024	Roth for Congress	$3,300, subsequently redesignated

The first two are supported by FEC filings; the Roth record includes a redesignation, so it shouldn't be counted as a separate additional donation without checking the transaction history. FEC.gov +2

A separate August 2026 investigation reports that she contributed **$1 million to Northwoods PAC**, supporting Jessi Ebben in Wisconsin's 7th Congressional District, as well as $7,000 directly to Ebben's campaign. The Badger Pr...

Search the complete FEC record

FEC individual contributions — Elizabeth Uihlein ↗

This is the official FEC search, where you can filter by election cycle, date, committee, and amount.

I also verified the Northwoods PAC's official FEC committee record ↗. FEC

OpenFEC skill status: I checked the skills available in this session, but your installed OpenFEC skill wasn't exposed. The records above come from public web searches and FEC filings, **not an execution of your skill**. Consequently, I haven't calculated a comprehensive federal contribution total or exported the complete transaction history.

The most useful next step would be to run the OpenFEC contributor search for `UIHLEIN, ELIZABETH`, including the `ELIZABETH A` name variant, and deduplicate amendments, redesignations, and memo transactions before aggregating contributions by recipient and election cycle.

Worked for 45s

I found several LinkedIn profiles for people named **Steven Johnson**, but there isn't enough information to determine which one you're researching.

LinkedIn matches

Person		
Author of *The Ghos* Google Labs		
Arts editor at The W		
Sports journalist, Fort		
Higher-education consultant, Mesa Park Strategies	Steven Johnson ↗	Austin
Business development and communications, San Diego Theatres	Steven Johnson ↗	San Diego

These are verified profile search matches, not evidence that any of them is the intended person. LinkedIn +4

One detail to narrow the search

Which Steven Johnson are you investigating? A city/state, employer, or approximate age would help me match the LinkedIn identity to FEC contributions without combining records from unrelated people.

FEC campaign contributions

I checked public FEC records and found at least one contribution reported under the name **Johnson, Steven**.

Field	Reported value
Na me	
Memo	Earmarked through WinRed

Source: FEC Schedule A filing ↗. I could not independently inspect the complete PDF, so this is based on the indexed filing text.

OpenFEC skill status: I checked the skills available in this conversation, and your custom OpenFEC skill isn't currently exposed to me. I therefore couldn't run the skill or perform an exhaustive contribution lookup.

You can also search the FEC individual contributions database ↗ directly.

“With this name and partial address, look this person up and find a little summary about them and grab their LinkedIn and details on any campaign contributions: {name} {address}”



STEP TWO:
Make it structured



Elizabeth Uihlein

Co-founder and president of Uline, the Wisconsin-based packaging and shipping supplies company. She and her husband, Richard Uihlein, are major contributors to Republican candidates and political committees. DeSmog +1

LinkedIn

I couldn't verify a personal LinkedIn profile for Elizabeth Uihlein. Searches for both "Elizabeth Uihlein" and "Liz Uihlein," including searches restricted to LinkedIn profiles, did not surface a credible match.

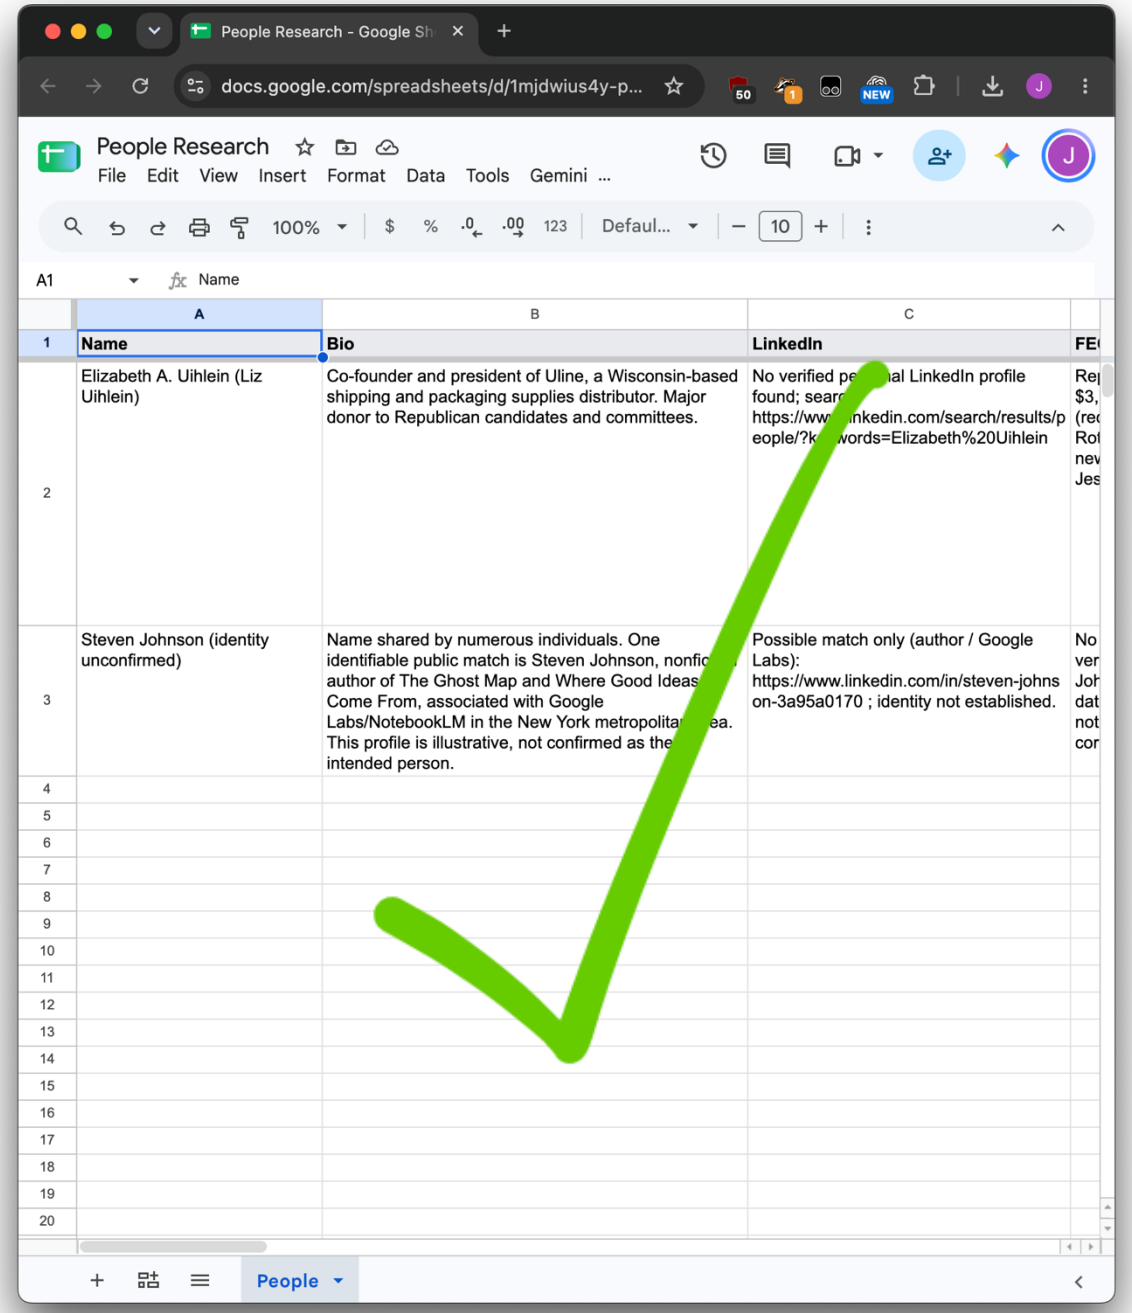
Search Elizabeth Uihlein on LinkedIn

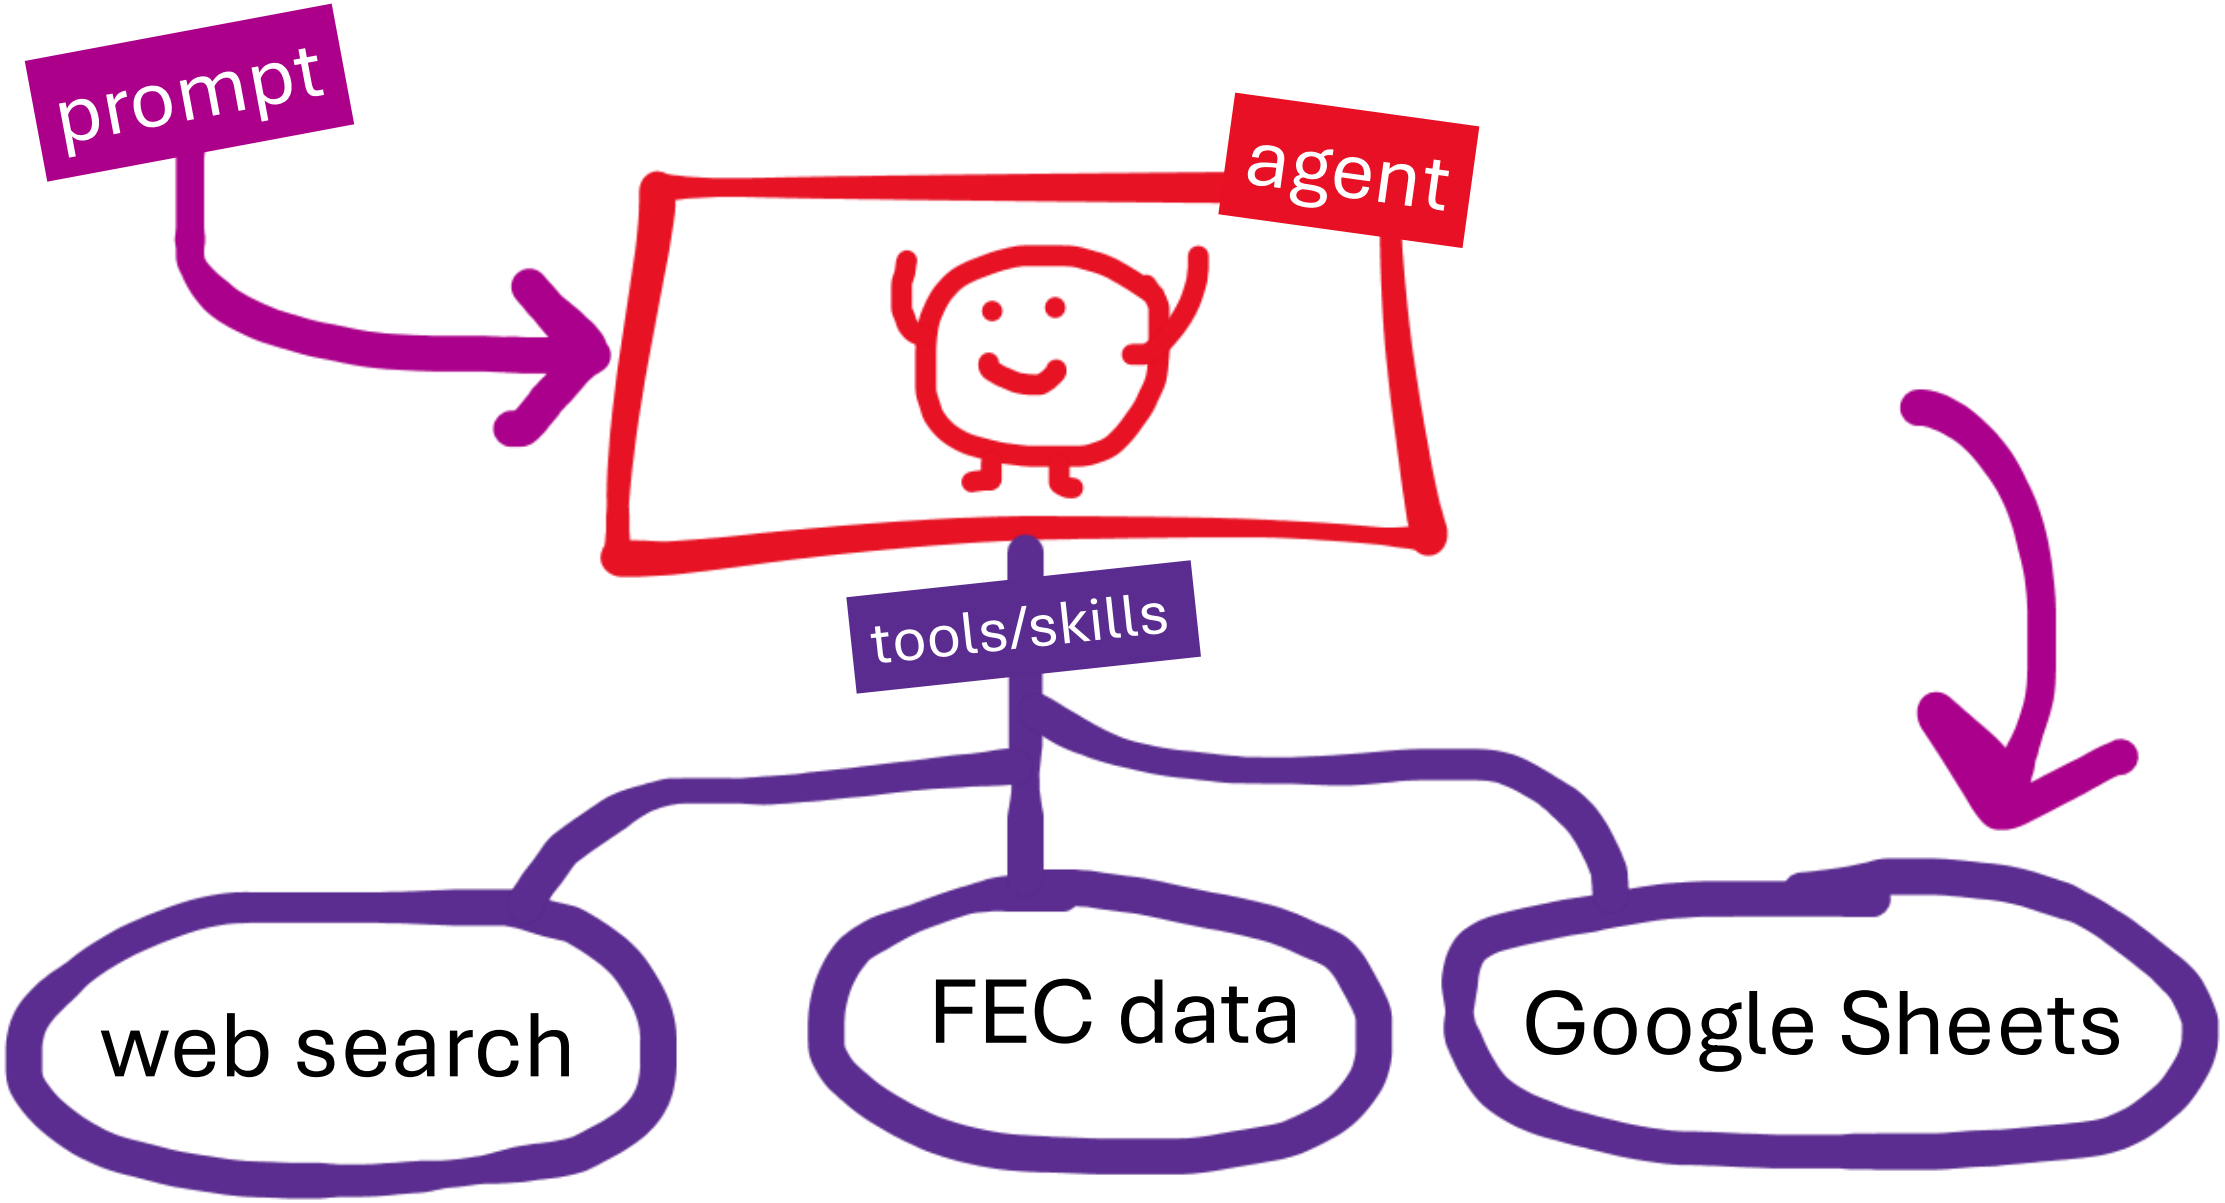
Federal campaign contributions

I found federal FEC filings identifying Elizabeth A. Uihlein as a contributor, often listing her employer as Uline and her occupation as CEO, president, or co-founder.

Here are some documented examples:

Contribution date	Recipient	Amount
May 21, 2026	Julia Letlow for Louisiana	$3,500 × 2
October 14, 2025	Federal candidate committee (filing excerpt)	$3,500
June 10, 2024	Roth for Congress	$3,300, subsequently redesignated

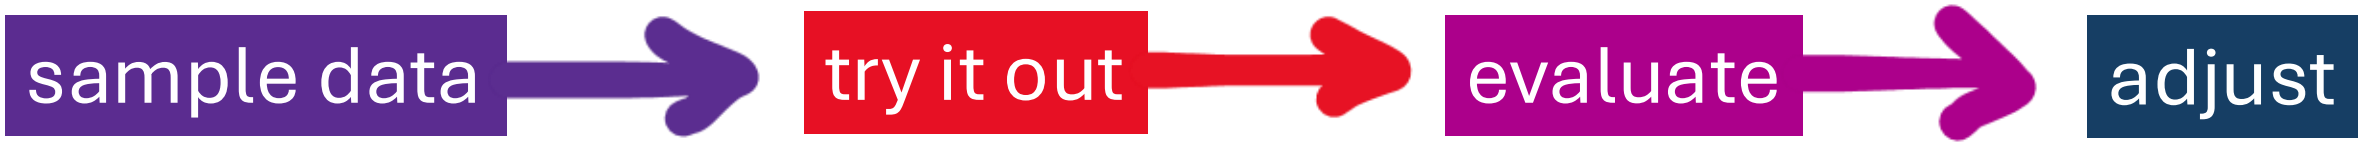



prompt
agent
tools/skills
web search
FEC data
Google Sheets

Structured places data can live

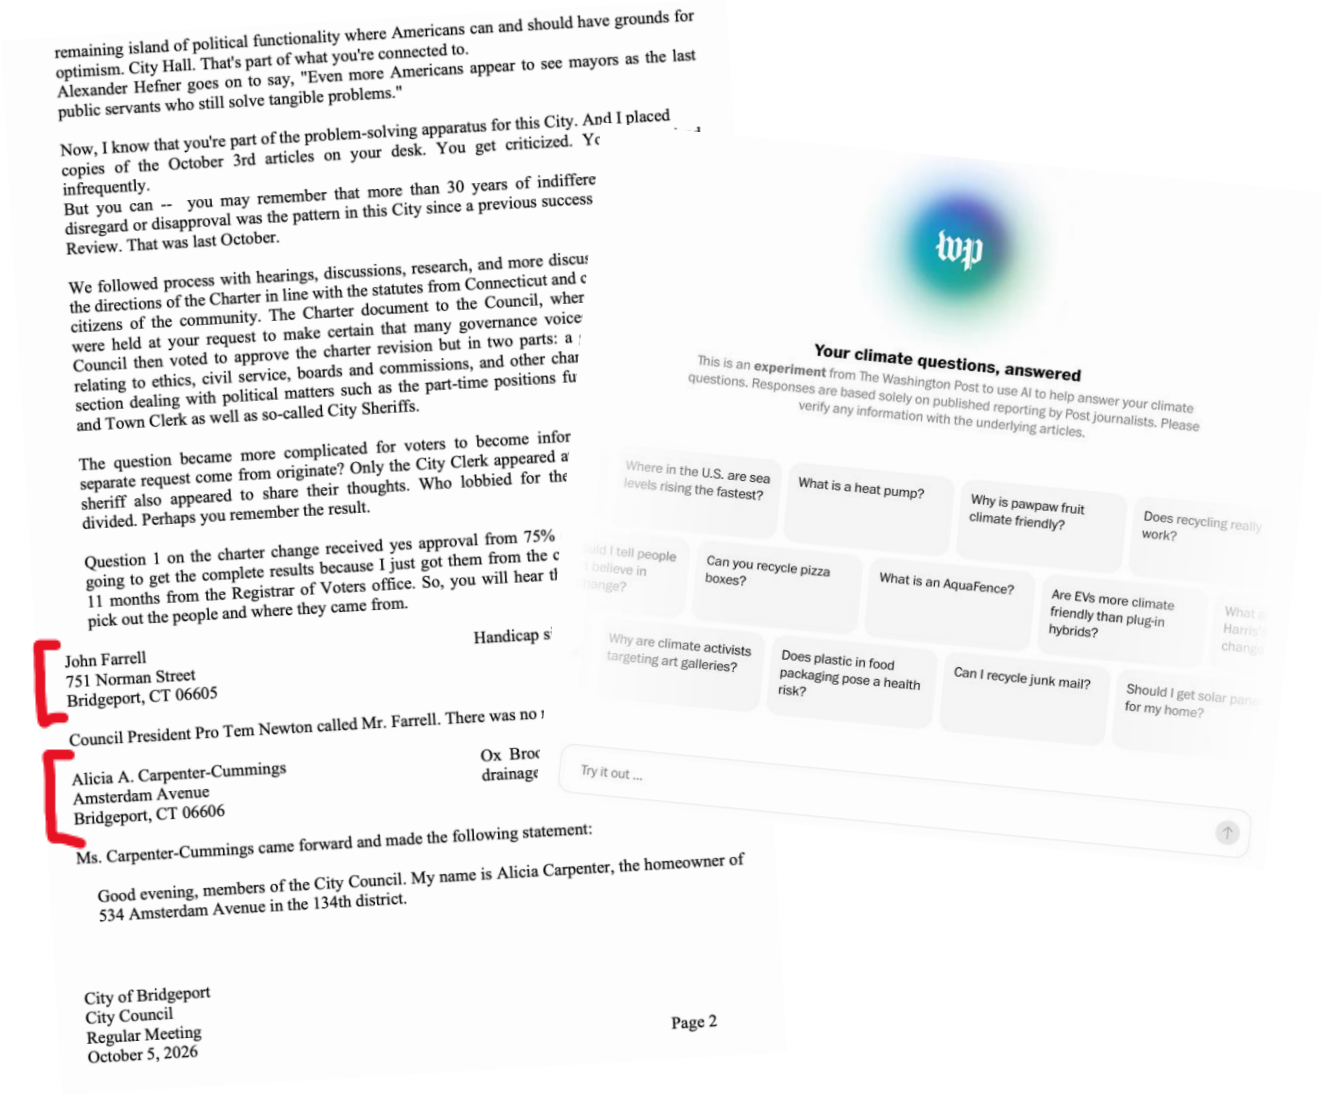
- Spreadsheets, obviously
- **Emails!** One email = one data point
- **Slack messages!** One post = one data point
- **GitHub!** One issue = one data point
- **Google Docs!** One file = one data point

STEP THREE:
Design and run your eval


sample data
try it out
evaluate
adjust

sample data → try it out → evaluate → adjust

- Real data is best!
- …but isn't always possible
- I love **synthetic data**





- ✕ “Hey ChatGPT, give me 100 sample data points”
- How could each person be ~~different~~ easier or more difficult to do research on, or somehow affect the results?
- **Dimensions of difference:** Which aspects have the most variance?
- Some options for doing research on people:
  - Rarity of name
  - Abbreviated name (Elizabeth vs Liz vs Beth)
  - Not online vs very online, has LinkedIn vs doesn’t
  - Full address vs partial address vs no address
- We’re researching real people so synthetic data is much much harder to make because it has to be *real people*. **It’s simpler to fake documents to analyze or reader questions to answer!**

Introducing the United States

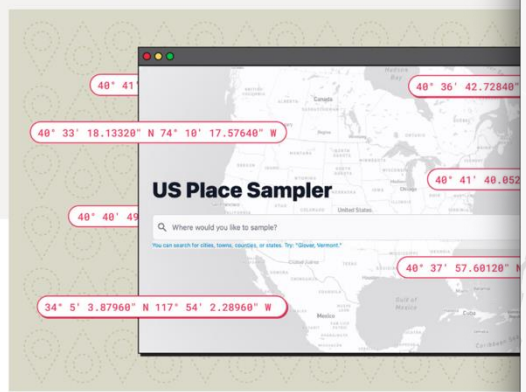
themarkup.org/still-loading/2023/03/02/introducing-the-united-states-place-sampler-usps

Challenging technology to serve the public good.

Still Loading

Introducing the United States Place Sampler (USPS)

Easily sample random U.S. street addresses using this ne
Local News and The Markup
By Leon Yin, Aaron Sankin, and Joe Nudell for Big Local N
Stanford
March 2, 2023 10:00 a.m. UTC



Gabriel Hongsdusit

Share This Article

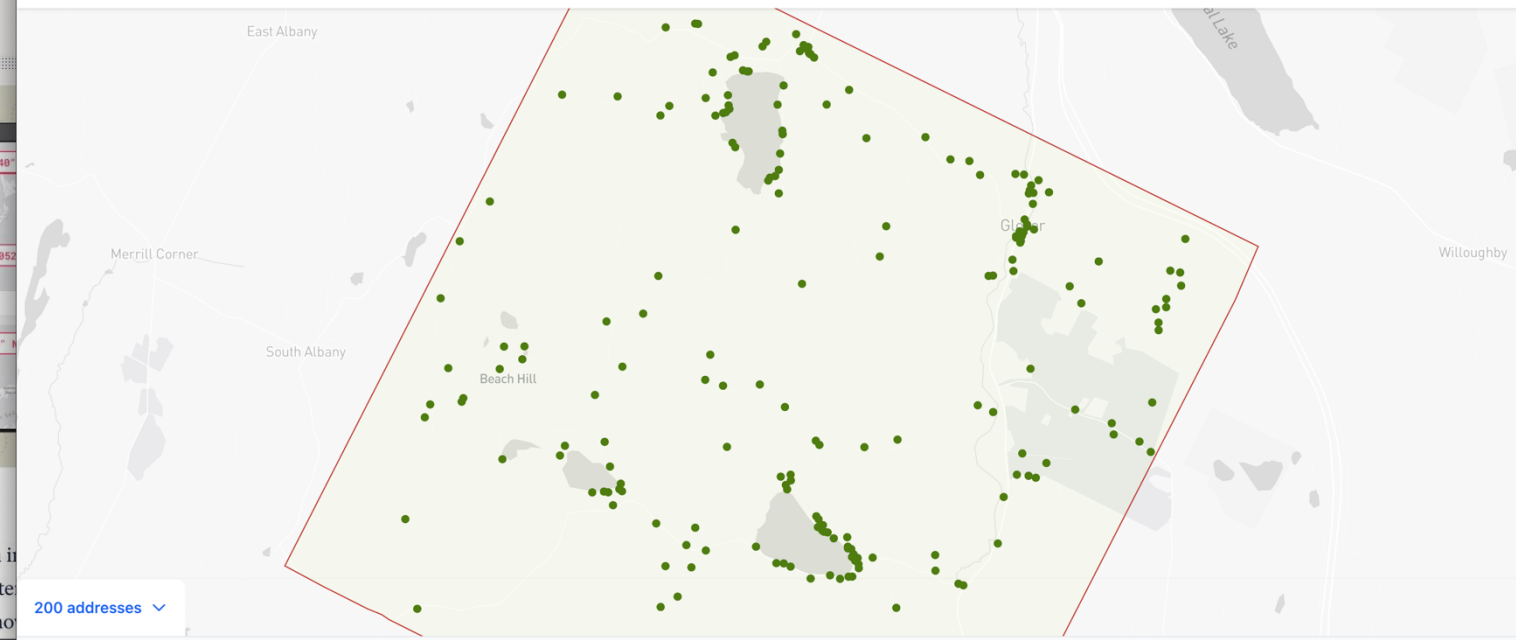
Copy Link Republish

In October 2022, The Markup published an investigation i
disparities in 38 major cities across the United States. Afte
readers, including fellow journalists, reached out to ask ho

Introducing the United States | U.S. Place Sampler: Sample a

usps.biglocalnews.org

U.S.P.S. I would like 200 total addresses from Glover × Sample



200 addresses

ADDRESS	COORDINATES
148 QUARRY HL GLOVER, VT 05839	44° 39′ 34.87680″ N 72° 12′ 32.58360″ W
735 DEXTER MOUNTAIN RD GLOVER, VT 05839	44° 41′ 58.40160″ N 72° 10′ 43.76640″ W
65 TRUCOTT PL GLOVER, VT 05839	44° 40′ 34.34880″ N 72° 11′ 16.89360″ W
2267 MUD ISLAND RD GLOVER, VT 05875	44° 39′ 34.48440″ N 72° 17′ 33.33840″ W



- Generate based on **combinations**
- All combinations:
  - (A1, B1, C1, D1), (A2, B1, C1, D1), (A2, B2, C1, D1)
  - Unusual name + no abbrevs + we have address + exists online
  - Unusual name + no abbrevs + we have address + doesn't exist online
  - Common name + yes abbrevs + no address + exists online
  - …etc…
- Use your intuition, filter what makes sense
  - If someone is lying we can't really check anything!
- End up with **100 total examples**

sample data → **try it out** → evaluate → adjust

- Have a single, trackable source of input data
- The easy/boring part: **do the work!**
- Depending on your process, cut+paste into new convo or “Hey ChatGPT, read the sheet and run the research skill on every one of the rows”
- You can run it multiple times for the same inputs, answers might vary!

Full name	Address
John Smith	New York
Michael Johnson	10001
Jennifer Williams	Chicago, IL
David Brown	
Maria Garcia	742 Evergreen Terrace, Springfield, IL
James Miller	California
Sarah Thompson	Brooklyn, NY 11215
Robert Martinez	2139
Emily Anderson	1847 Maple Avenue, Denver, CO 80203
Daniel Wilson	Austin, TX
Rebecca Sullivan	Florida
Marcus Bennett	
Caroline Fletcher	60614
Anthony Delgado	2901 Oak Street, Seattle, WA
Priya Shah	Philadelphia, PA 19103
Nathaniel Brooks	New Jersey
Elena Whitcomb	33139
Dominic Castellano	563 Willow Lane, Portland, OR 97205
Vivian Holloway	Boston, MA
Malcolm Prescott	
Jasper Wainwright	90210
Amara Okonkwo	Minneapolis, MN 55401
Celeste Vandenberg	Georgia



- **Golden** or **ground truth dataset:** It has the answers, did the actual work!
- Keep it in a **separate sheet** so the LLM doesn't cheat

A	B	C	D	E
Full name	**Address**	**LinkedIn**	**Bio**	**Contributions**
John Smith	New York	https://example.com/linkedin/joh	Software engineer in New York; employer unverified	Multiple John Smith records; identity unclear
Michael Johnson	10001	Not found	Former sales manager, possibly based in Chicago	Not found
Jennifer Williams	Chicago, IL	https://example.com/linkedin/jen	Marketing director at a technology company in Denver	$500 to Biden for President (2020); $250 to DNC (2022)
David Brown		Multiple possible matches	Not found	$1,000 to Trump for President (2020); identity uncertain
Maria Garcia	742 Evergreen Terrace, S	https://example.com/linkedin/ma	Community organizer in Los Angeles	Not found
James Miller	California	https://example.com/linkedin/jan	Financial analyst at a Chicago investment firm	$2,800 to Biden for President (2020)
			Elementary school teacher in	

sample data → try it out → **evaluate** → adjust

- DO THIS MANUALLY.
- *DO THIS MANUALLY.*
- ***DO THIS MANUALLY.***
- Only pass/fail binary!! **No scores!!!!!**
- Helps you define the actual boundary
- Multiple reviewers – research **“inter-rater reliability”** to be fancy

Success?	Note
Yes	
No	"Elizabeth" did not match with "Liz"
No	Used middle name when filing paperwork
No	"Jon" did not match with "John"
Yes	
Yes	
No	Wrong person, claimed NYC and New Jersey were "close enough"
Yes	
Yes	
Yes	
	Claimed uncertainty about match even

sample data → try it out → **evaluate** → adjust

- …you can have an LLM do this if you first compare its results to your results and find out it agrees with you.
- “Given the prompt and the results below, did the model meet the following criteria? <X, Y, Z>”
- Weak/cheap model for the actual process, **strong model as reviewer**

Success?	Note
Yes	
No	"Elizabeth" did not match with "Liz"
No	Used middle name when filing paperwork
No	"Jon" did not match with "John"
Yes	
Yes	
No	Wrong person, claimed NYC and New Jersey were "close enough"
Yes	
Yes	
Yes	
	Claimed uncertainty about match even



- Read your free-form error summaries and **categorize your problems**
- Fix them systematically

```
prompt = """
Look this person up and find a little summary
about them and grab their LinkedIn and details on
any campaign contributions.

- If a name has diminutives, perform multiple
searches e.g. Elizabeth->Liz/Beth/Eliza/Bess
- If there is a first + middle name, search with
it, without it, and with it abbreviated
- Err on the side of making assumptions, your
findings will be fact-checked.
- People listed as being in NYC might be in New
Jersey or Connecticut
- If you're given the ZIP code, use that to look
up the city/state first
- Skip common names if you don't have any address
information
"""
```

sample data → try it out → evaluate → adjust

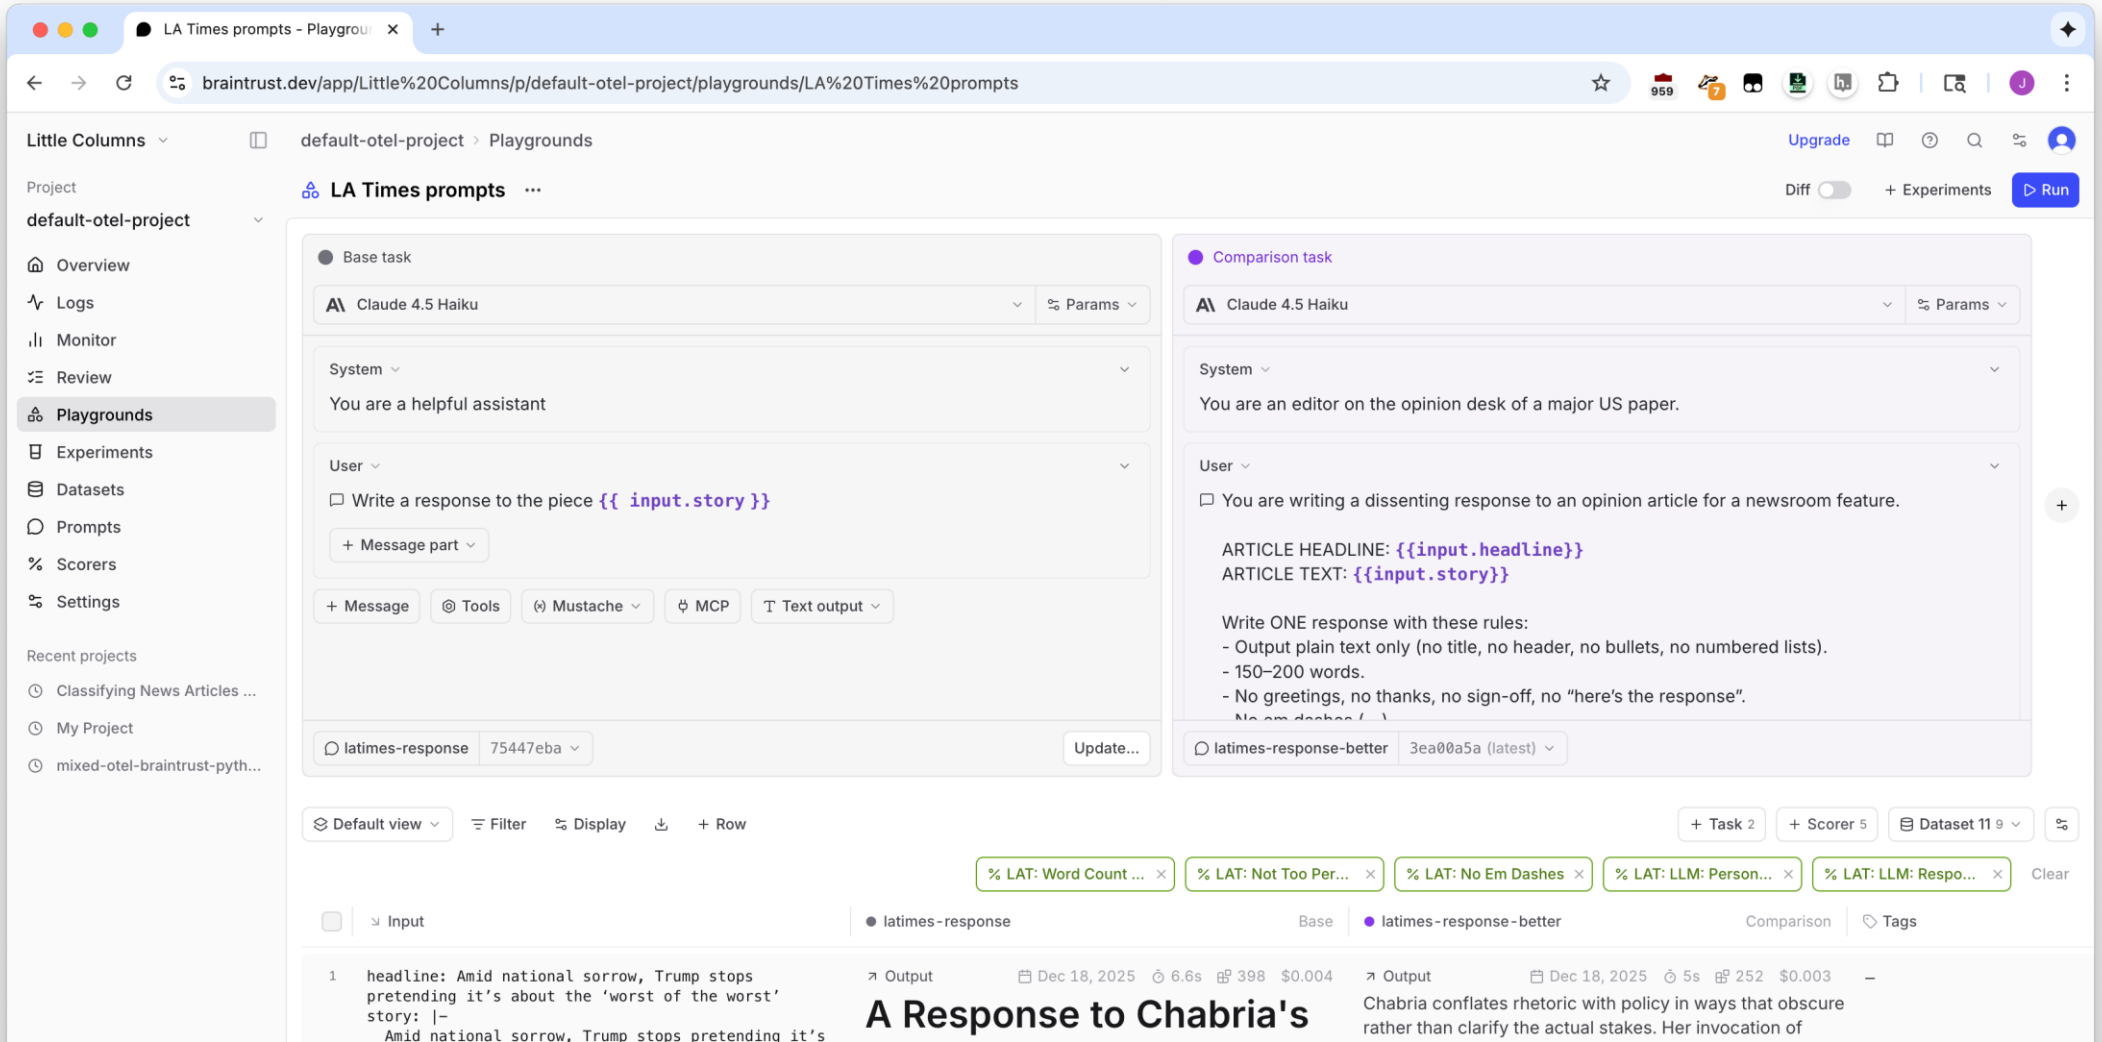
- Examples: “if X, do Y”
- I am sorry but ***YOU CAN’T RE-USE YOUR SAMPLES FOR THE NEXT EVAL.***
  - You’d be teaching to the test!
- **Train/test split:** keep some hidden from whoever is adjusting the prompt to do the evaluation

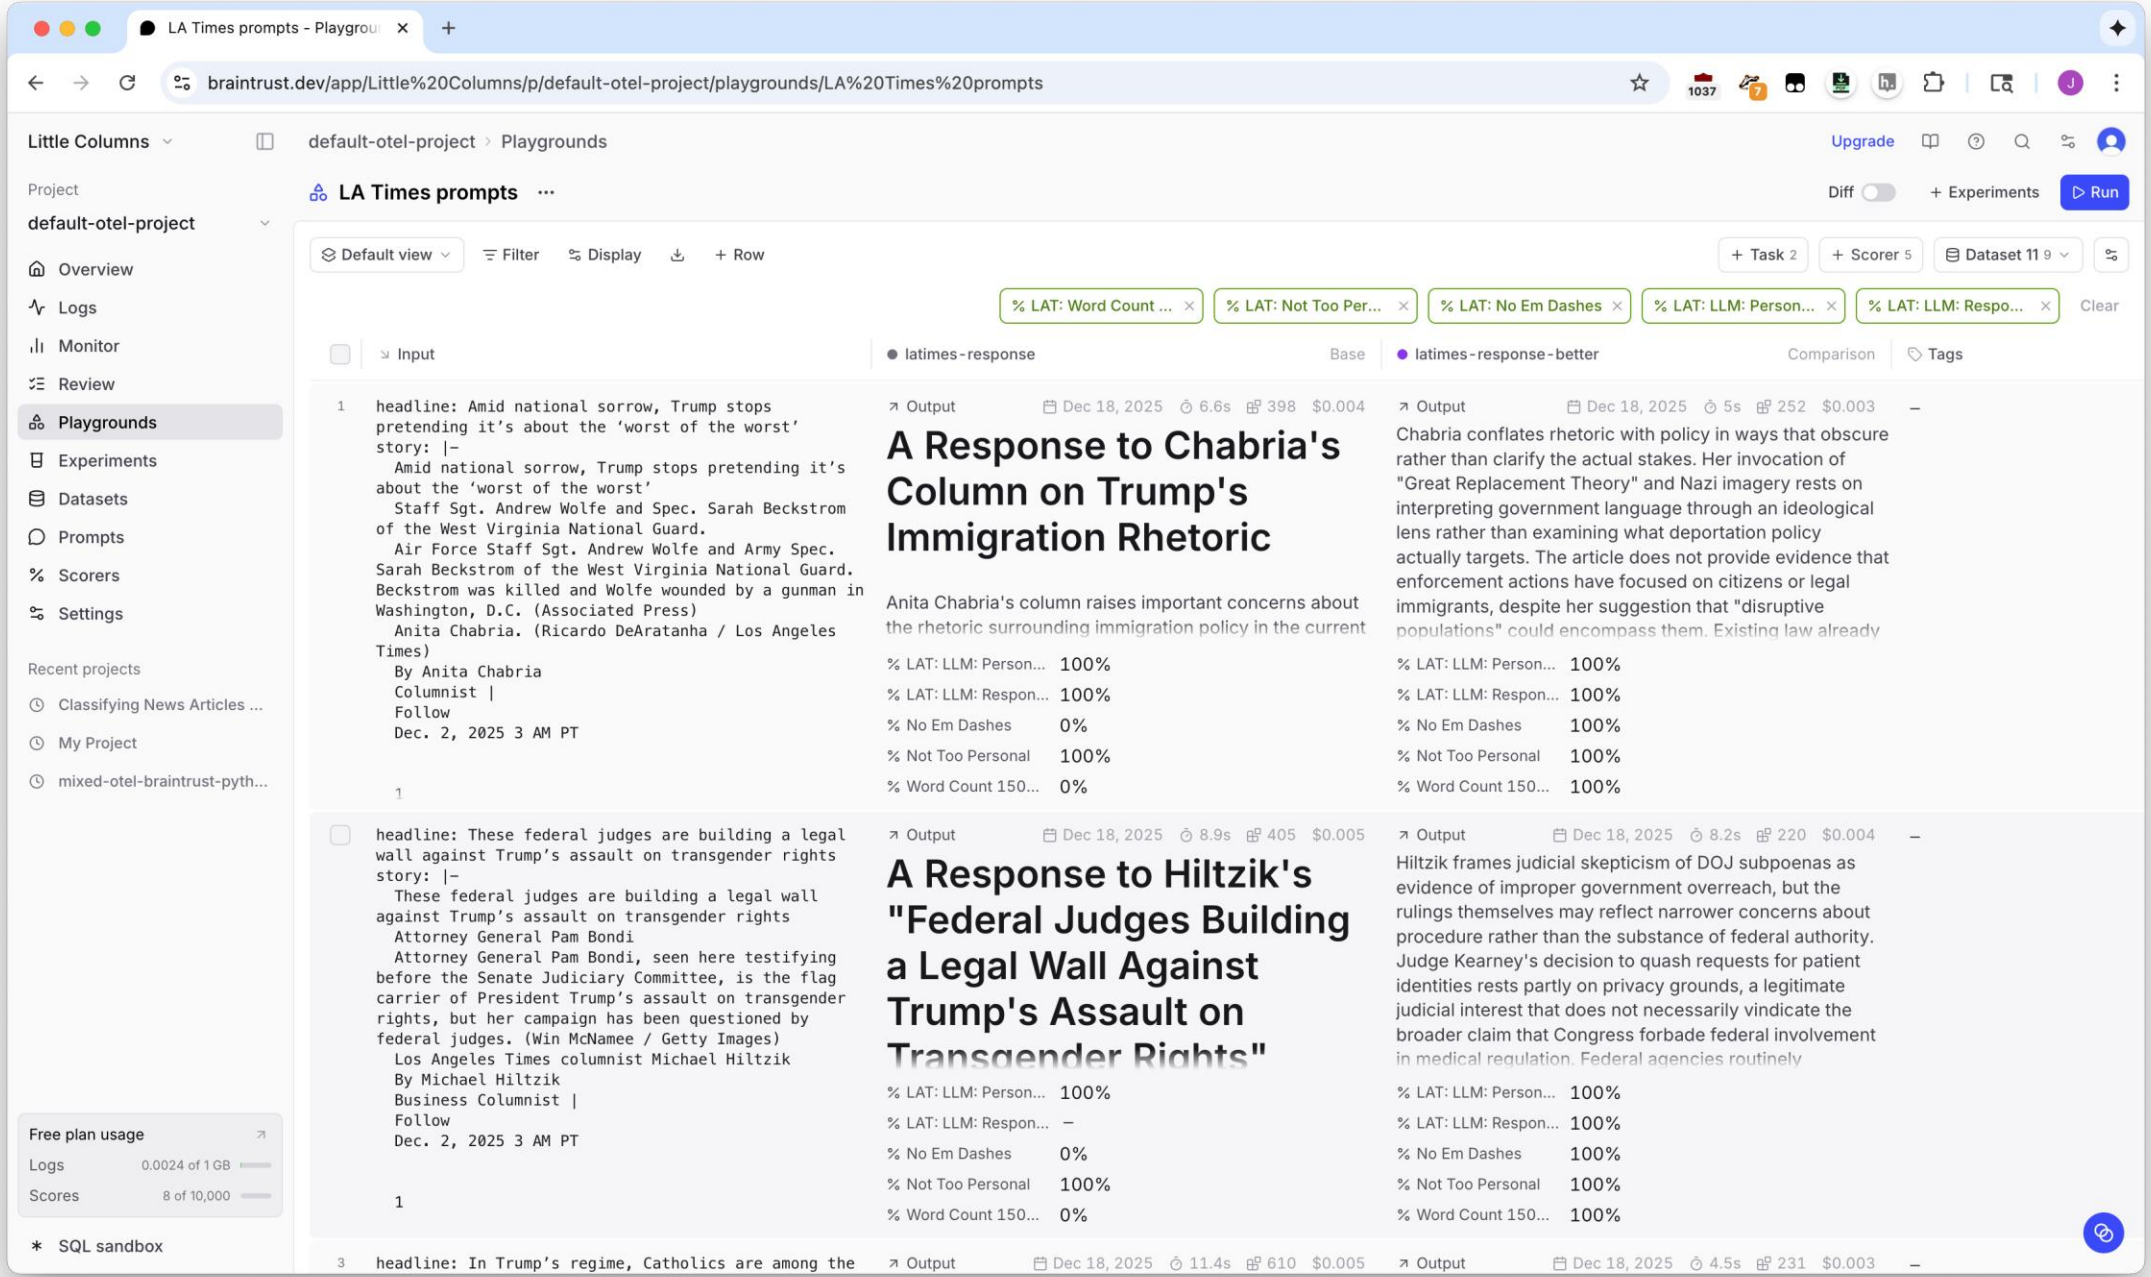
```
prompt = """
Look this person up and find a little summary
about them and grab their LinkedIn and details on
any campaign contributions.

- If a name has diminutives, perform multiple
searches e.g. Elizabeth->Liz/Beth/Eliza/Bess
- If there is a first + middle name, search with
it, without it, and with it abbreviated
- Err on the side of making assumptions, your
findings will be fact-checked.
- People listed as being in NYC might be in New
Jersey or Connecticut
- If you're given the ZIP code, use that to look
up the city/state first
- Skip common names if you don't have any address
information
"""
```

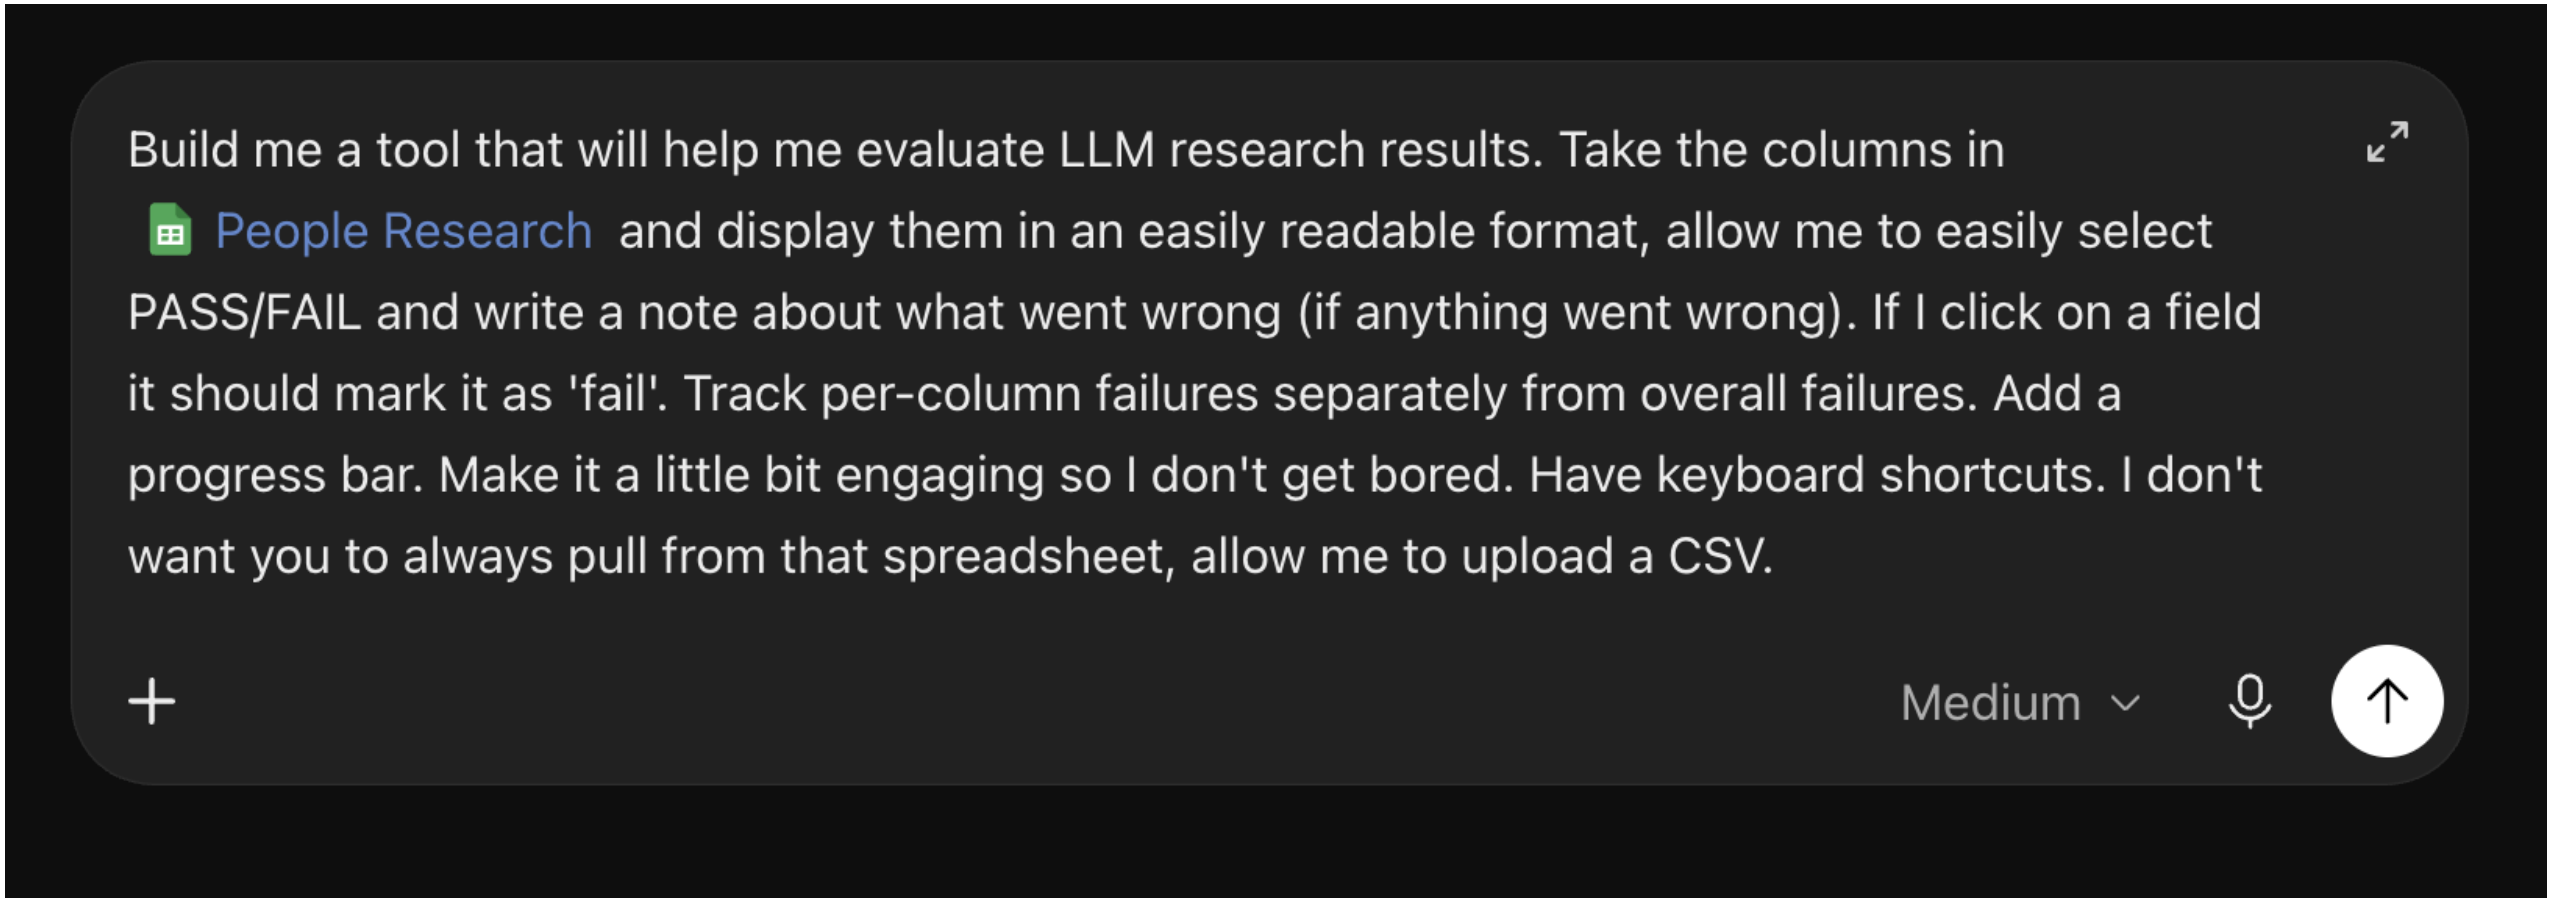
So how do we *actually* do this?

Eval tools: Braintrust, Langfuse, etc




LA Times prompts - Playgrou
braintrust.dev/app/Little%20Columns/p/default-otel-project/playgrounds/LA%20Times%20prompts
Little Columns
default-otel-project › Playgrounds
Upgrade
Project
default-otel-project
Overview
Logs
Monitor
Review
Playgrounds
Experiments
Datasets
Prompts
Scorers
Settings
Recent projects
Classifying News Articles ...
My Project
mixed-otel-braintrust-pyth...
LA Times prompts
Diff
+ Experiments
Run
Default view
Filter
Display
+ Row
+ Task 2
+ Scorer 5
Dataset 11 9
% LAT: Word Count ...
% LAT: Not Too Per...
% LAT: No Em Dashes
% LAT: LLM: Person...
% LAT: LLM: Respo...
Clear
Input
latimes-response
Base
latimes-response-better
Comparison
Tags
1
headline: Amid national sorrow, Trump stops pretending it's about the 'worst of the worst' story: |-
Amid national sorrow, Trump stops pretending it's about the 'worst of the worst'
Staff Sgt. Andrew Wolfe and Spec. Sarah Beckstrom of the West Virginia National Guard.
Air Force Staff Sgt. Andrew Wolfe and Army Spec. Sarah Beckstrom of the West Virginia National Guard. Beckstrom was killed and Wolfe wounded by a gunman in Washington, D.C. (Associated Press)
Anita Chabria. (Ricardo DeAratanha / Los Angeles Times)
By Anita Chabria
Columnist |
Follow
Dec. 2, 2025 3 AM PT
1
Output
Dec 18, 2025
6.6s
398
$0.004
A Response to Chabria's Column on Trump's Immigration Rhetoric
Anita Chabria's column raises important concerns about the rhetoric surrounding immigration policy in the current
% LAT: LLM: Person... 100%
% LAT: LLM: Respon... 100%
% No Em Dashes 0%
% Not Too Personal 100%
% Word Count 150... 0%
Output
Dec 18, 2025
5s
252
$0.003
Chabria conflates rhetoric with policy in ways that obscure rather than clarify the actual stakes. Her invocation of "Great Replacement Theory" and Nazi imagery rests on interpreting government language through an ideological lens rather than examining what deportation policy actually targets. The article does not provide evidence that enforcement actions have focused on citizens or legal immigrants, despite her suggestion that "disruptive populations" could encompass them. Existing law already
% LAT: LLM: Person... 100%
% LAT: LLM: Respon... 100%
% No Em Dashes 100%
% Not Too Personal 100%
% Word Count 150... 100%
headline: These federal judges are building a legal wall against Trump's assault on transgender rights story: |-
These federal judges are building a legal wall against Trump's assault on transgender rights
Attorney General Pam Bondi
Attorney General Pam Bondi, seen here testifying before the Senate Judiciary Committee, is the flag carrier of President Trump's assault on transgender rights, but her campaign has been questioned by federal judges. (Win McNamee / Getty Images)
Los Angeles Times columnist Michael Hiltzik
By Michael Hiltzik
Business Columnist |
Follow
Dec. 2, 2025 3 AM PT
1
Output
Dec 18, 2025
8.9s
405
$0.005
A Response to Hiltzik's "Federal Judges Building a Legal Wall Against Trump's Assault on Transgender Rights"
% LAT: LLM: Person... 100%
% LAT: LLM: Respon... –
% No Em Dashes 0%
% Not Too Personal 100%
% Word Count 150... 0%
Output
Dec 18, 2025
8.2s
220
$0.004
Hiltzik frames judicial skepticism of DOJ subpoenas as evidence of improper government overreach, but the rulings themselves may reflect narrower concerns about procedure rather than the substance of federal authority. Judge Kearney's decision to quash requests for patient identities rests partly on privacy grounds, a legitimate judicial interest that does not necessarily vindicate the broader claim that Congress forbade federal involvement in medical regulation. Federal agencies routinely
% LAT: LLM: Person... 100%
% LAT: LLM: Respon... 100%
% No Em Dashes 100%
% Not Too Personal 100%
% Word Count 150... 100%
3
headline: In Trump's regime, Catholics are among the
Output
Dec 18, 2025
11.4s
610
$0.005
Output
Dec 18, 2025
4.5s
231
$0.003
Free plan usage
Logs 0.0024 of 1 GB
Scores 8 of 10,000
SQL sandbox

Software is free! Just build a custom tool



Research Review

People Research - People.csv · 2 people · Autosaved locally

Import CSV | Reset | Export CSV | Next →

Progress
1 / 2
50% reviewed · 1 remaining

Overall passed	Overall failed	Field failures	Streak
1	0	6	1

Review queue 2

All records

- **Elizabeth A. Uihlein (Liz Uihlein)**
  #1 · PASS · 3 flagged
- **Steven Johnson (identity unconfirmed)**
  #2 · Unreviewed · 3 flagged

RECORD 2 OF 2

Steven Johnson (identity unconfirmed)

Awaiting overall verdict

3 fields flagged

Research fields

1. Name	Steven Johnson (identity unconfirmed)
2. Bio FAIL	Name shared by numerous individuals. One identifiable public match is Steven Johnson, nonfiction author of The Ghost Map and Where Good Ideas Come From, associated with Google Labs/NotebookLM in the New York metropolitan area. This profile is illustrative, not confirmed as the intended person.
3. LinkedIn	Possible match only (author / Google Labs): https://www.linkedin.com/in/steven-johnson-3a95a0170 ; identity not established.
4. FEC campaign contributions FAIL	No contributions attributable to the intended Steven Johnson have been verified. A previously surfaced FEC filing was described as listing a Steven Johnson in Cheverly, MD, self-employed engineer, with a $25 contribution dated Sept. 13, 2024; this is a DIFFERENT/UNLINKED identity candidate, not attributable to the LinkedIn profile or intended subject without further corroboration.
5. Contribution sources FAIL	https://www.fec.gov/data/individual-contributions/?contributor_name=Steven+Johnson ; https://docquery.fec.gov/pdf/208/202504029755014208/202504029755014208.pdf (filing detail unverified)
6. Bio / LinkedIn sources	https://www.linkedin.com/in/steven-johnson-3a95a0170
7. Research notes	High identity ambiguity: many people called Steven Johnson. Author/Google Labs LinkedIn profile is genuine but not confirmed as the subject. Do not associate FEC candidate with that profile; employer/location differ. Custom OpenFEC skill unavailable in this conversation; no exhaustive FEC query. Added Oct. 20, 2026.

Click a row to mark a field as failed; click again to clear. Overall verdict is separate.

Overall verdict

✓ PASS [P] | ✕ FAIL [F] | ↱ SKIP [S]

Notes

What failed? Wrong identity, hallucinated sources, incomplete contributions, unclear evidence...

← Previous | Save & next →

Field failures

Per-column failures across the entire dataset.

Name	1
Bio	2
LinkedIn	0
FEC campaign contributions	2
Contribution sources	1
Bio / LinkedIn sources	0
Research notes	0

Keyboard controls

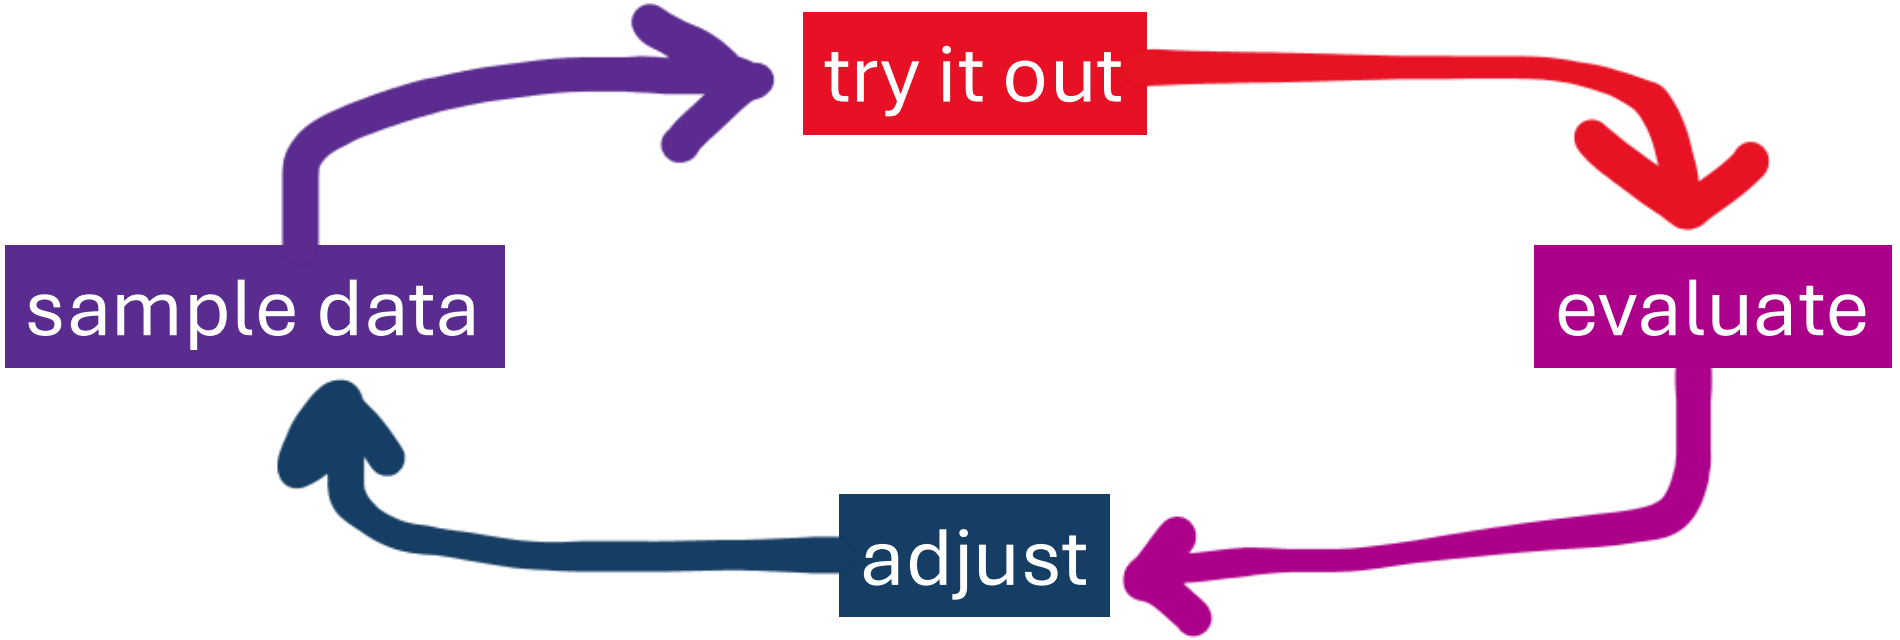
- P Overall pass
- F Overall fail
- S Skip
- → Next record
- ← Previous record
- 1–9 Toggle field failure
- N Focus notes
- Esc Unfocus notes

Keyboard shortcuts are suspended while typing in notes.

Reviews are saved in this browser. Export CSV to back up your work.

What to do with outputs

- Tweak the prompts
  - Examples are great!
- Prompt engineering is dead, long live prompt engineering



Advanced work: Peering in the box

- **Traces:** Everything that happens
- **Tool calls:** Web search, etc
- **Instrumentation**: Connecting your AI to measuring tools
- **Observability:** Using instrumentation to track and analyze
- **Telemetry:** A more generic word for instrumentation

```
from openai import OpenAI
from somewhere import openfec_data, exa

client = OpenAI()

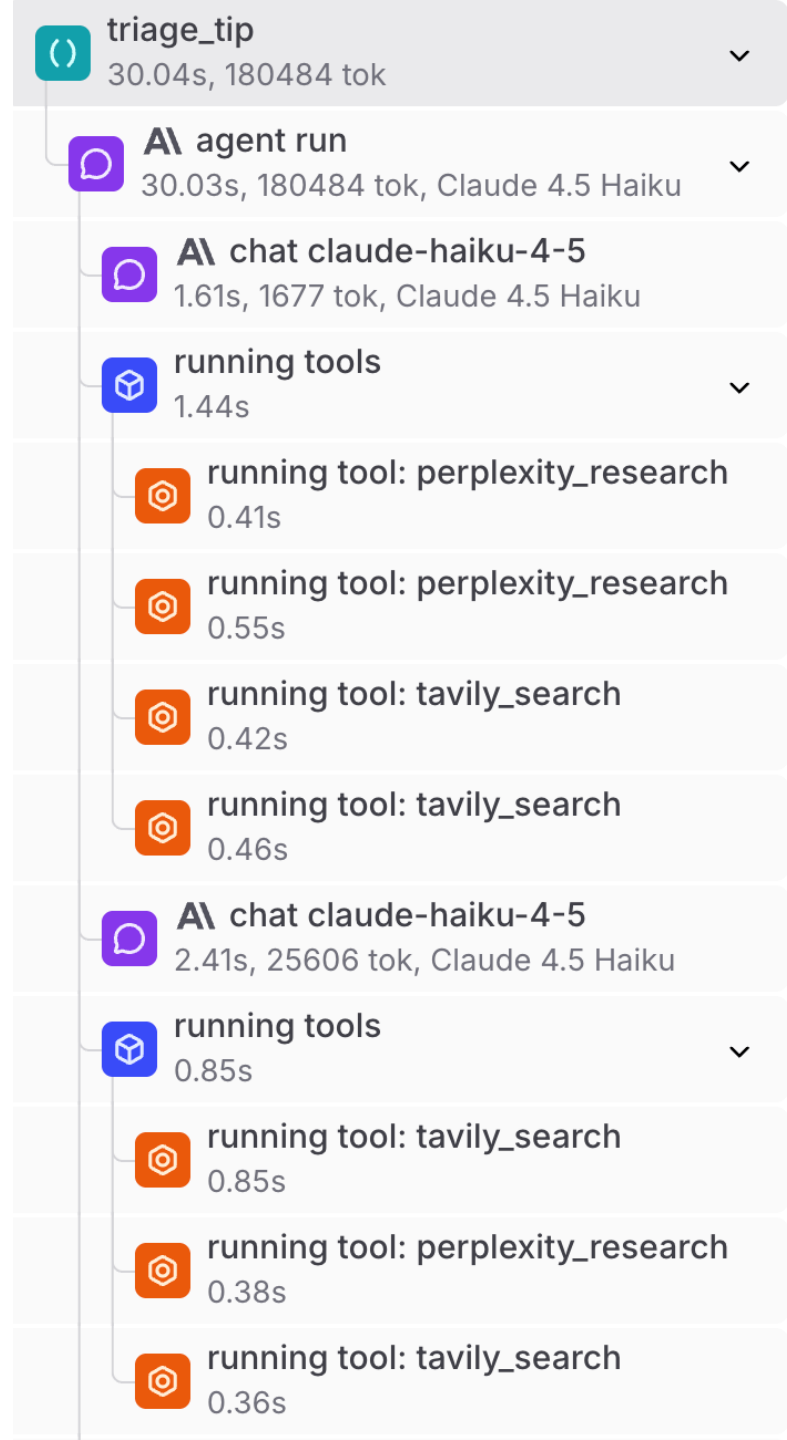
# Imaginary helper converts Python func
into API tool definitions.
response = client.responses.create(
    model="gpt-5",
    input="Research {name}'s businesses
political donations.",
    tools=[
        tool(exa_search),
        tool(openfec_data),
    ],
)
```


import OpenAI from "openai";
import { initLogger } from "braintrust";
initLogger({
projectName: "My Project", // Replace with your project name
apiKey: process.env.BRAINTRUST_API_KEY,
});
const client = new OpenAI({ apiKey: process.env.OPENAI_API_KEY });
const result = await client.chat.completions.create({
model: "gpt-5-mini",
messages: [{ role: "user", content: "What is machine learning?" }],
});
Braintrust added

```
You are an investigative journalism assistant that triages
news tips.

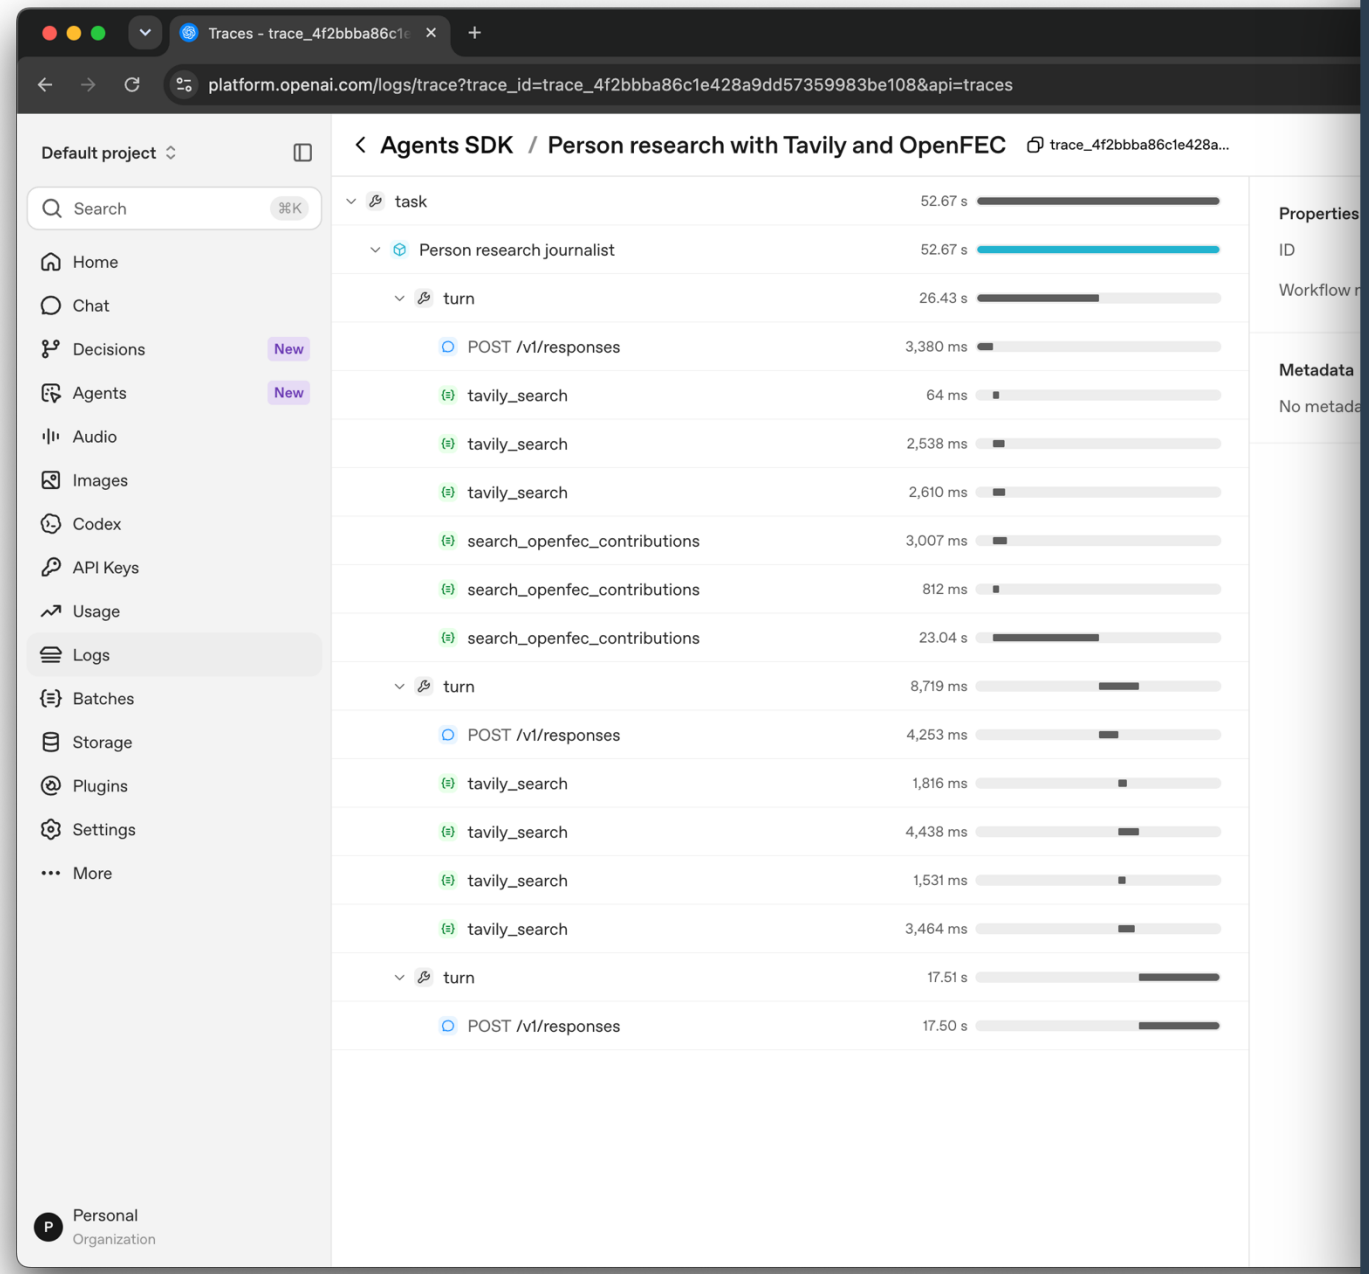
Your task is to:
1. Evaluate the tip for credibility and newsworthiness (score
1-10)
2. Determine if it is valid (score >= 5) or spam/invalid
(score < 5)
3. Categorize it into the most appropriate beat from this
list:
4. Write a concise 2-3 sentence summary
5. Use Perplexity search to conduct background research on the
topic, entities, and location mentioned
6. Research the tipster's name and email to assess credibility
(look for public profiles, previous journalism involvement,
potential biases)
7. Generate 3-5 specific follow-up questions for further
investigation
8. Compose a professional email response:
   - If valid: Thank them, indicate you're reviewing it, and
ask the follow-up questions
   - If invalid: Politely decline and thank them for their
interest

Be thorough in your research but efficient. Focus on
verifiable facts and credible sources. Understand that people
don't really know how to write tips, and you should follow up
on most of the ones that are not obviously crazy people.
```



Just vibe-code yourself into traces

Write me some sample code that uses the agents sdk. It should use the openfec skill that we made along with web search to research a person. Give me a downloadable package that I can run if i add an openai api key




```python
def build_agent(model: str) -> Agent:
    skill = (ROOT / "skills/openfec/SKILL.md"
    return Agent(
        name="Person research journalist",
        model=model,
        instructions=(
            "Research a person for a journali
            "Use the tavily_search tool (not
            "profile if one exists, and sourc
            "and variants when useful. Clearl
            "matches. Never invent a LinkedIn
            "and the official FEC data source
            "Output a concise Markdown resear
            "records (date, amount, committee
            "caveats and next verification st
            "OPENFEC SKILL INSTRUCTIONS:\n" +
        ),
        tools=[tavily_search, search_openfec_
    )
```

task 52.67 s

Person research journalist 52.67 s

turn 26.43 s

POST /v1/responses 3,380 ms

tavily_search 64 ms

tavily_search 2,538 ms

Function call

```
search_openfec_contributions({
  "name": "Elizabeth Uihlein",
  "state": "WI",
  "city": null,
  "min_date": null,
  "max_date": null,
  "max_pages": 2
})
```

turn 17.51 s

POST /v1/responses 17.50 s

tavily

Prop

Crea

Func

Function call

```
search_oper
  "name": "
  "state":
  "city": n
  "min_date
  "max_date
  "max_page
})
```

Output

```
{
  "query":
    "contri
    "contri
  },
  "source"
  "results"
  "results_
  "pages_fe
  "possibly
  "identity
}
```

Function call

```
search_openfec_contributions({
  "name": "Elizabeth A Uihlein",
  "state": "WI",
  "city": null,
  "min_date": null,
  "max_date": null,
  "max_pages": 2
})
```

Output


```
{
  "query": {
    "contributor_name": "Elizabeth A Uihlein",
    "contributor_state": "WI"
  },
  "source": "https://api.open.fec.gov/developers/",
  "results": [
    {
      "contributor_name": "UIHLEIN, ELIZABETH A.",
      "contributor_city": "PLEASANT PRAIRIE",
      "contributor_state": "WI",
      "contributor_employer": "ULINE",
      "contributor_occupation": "PRESIDENT",
      "contribution_receipt_date": "2026-06-02",
      "contribution_receipt_amount": 3500,
      "committee_id": "C00895078",
      "committee_name": "SMITH FOR MI-08 REPUBLICAN NOMINEE
FUND 2026",
      "sub_id": "4090920261596024719",
      "file_number": 2009578,
      "memo_code": "X",
```

A gotcha: **the bitter lesson**

A better model is just about to solve the problem without any extra prompting work

Evaluate: Story-to-blurb summarizer

- We are designing something to make it easy for our reporters to write blurbs for newsletters and/or social media. **It takes a story, outputs a blurb.**
- **What is our evaluation criteria?**
  - Honestly, we probably won't know what's good/bad until we iterate a bunch!!
- **What are our dimensions of difference?**
  - Beat/topic
  - Which social media channel it is going to/target newsletter audience
  - Level of seriousness
  - Probably 1,000 things we can't think of, this is how everyone gets bitten
- **How do we generate synthetic examples of entire stories?**
  - We don't! Just take the past year of examples
- **What does our spreadsheet look like?**
  - Do we need a gold set? If yes, how do we build it?
- **Do we need tracing?**

Evaluate: “Ask our archives” bot

- We’re designing a chatbot to answer questions about local development.
- **What is our evaluation criteria?**
  - Accurate: answer reflects reporting
  - Evidence-backed: citations support claims
  - More recent information supercedes older information
  - Acknowledges gaps
- **What are our dimensions of difference?**
  - Unrelated topic, red teaming, biased question, about old data, etc
- **What does our spreadsheet look like?**
  - Do we need a gold set? If yes, how do we build it?
- **Do we need tracing?**
  - Most likely! It will be important to dig into *how* it’s actually searching.

Stress-Testing AI:
How to Evaluate Whether Your AI Tools Are Working

Jonathan Soma
js4571@columbia.edu @dangerscarf
Columbia University